

Title: Automatic Differentiation with pytorch

Author: [Eugene Han](#) and Ruoqing Zhu

For some basic reading about how pytorch implements automatic differentiation at a high-level, you can look at this [link](#). This [paper](#) details why machine learning methods primarily uses automatic differentiation. This [paper](#) is a review of how it works and is where the table below is from.

Technique	Advantage(s)	Drawback(s)
Hand-coded analytical derivative	Exact and often fastest method.	Time consuming to code, error prone, and not applicable to problems with implicit solutions. Not automated.
Finite differentiation	Easy to code.	Subject to floating point precision errors and slow, especially in high dimensions, as the method requires at least D evaluations, where D is the number of partial derivatives required.
Symbolic differentiation	Exact, generates symbolic expressions.	Memory intensive and slow. Cannot handle statements such as unbounded loops.
Automatic differentiation	Exact, speed is comparable to hand-coding derivatives, highly applicable.	Needs to be carefully implemented, although this is already done in several packages.

Table 1: Summary of techniques to calculate derivatives.

Example

Let

$$f(x_1, x_2) = (\log(x_1 x_2), \sin(x_2)),$$

and

$$g(x_1, x_2) = x_1 x_2,$$

we want to find the gradients

```
In [ ]: import torch

def f(x):
    return((torch.log(x[0] * x[1]), torch.sin(x[1])))

def g(x):
    return(x[0] * x[1])
```

```
# computational graph is not recorded by default, need to set the requires_grad
x = torch.tensor([0.5, 0.75], requires_grad=True)
y = g(f(x))

# computes the gradient of the current tensor w.r.t graph leaves
y.backward()

print(x.grad)
```

```
tensor([1.3633, 0.1912])
```

Manual Derivative Check

Note that

$$\nabla g(f(x_1, x_2)) = \begin{bmatrix} \frac{\sin(x_2)}{x_1} \\ \frac{\sin(x_2)}{x_2} + \log(x_1 x_2) \cos(x_2) \end{bmatrix}$$

```
In [ ]: dx1 = (torch.sin(x[1])/x[0]).item()
dx2 = (torch.sin(x[1])/x[1] + torch.log(x[0]*x[1])*torch.cos(x[1])).item()
print(round(dx1, 4), round(dx2, 4))
```

```
1.3633 0.1912
```

Numerical Differentiation (Finite differences)

```
In [ ]: dx = torch.zeros(2)
dx[0] = (g(f(x + torch.tensor([1e-3, 0]))) - y)/1e-3
dx[1] = (g(f(x + torch.tensor([0, 1e-3]))) - y)/1e-3
print(dx)
```

```
tensor([1.3619, 0.1919], grad_fn=<CopySlices>)
```

Example: Use case in linear regression

Let's do this for a simple linear regression problem:

```
In [ ]: import torch
from torch.nn import Parameter # Import the Parameter class

# The data function is: y = x + 10
n = 100
x_train = torch.randn(n, dtype=torch.float)
y_train = 2 + x_train + torch.randn(n, dtype=torch.float)

# Initialize Linear Regression parameters
b0 = torch.randn(1, requires_grad=True, dtype=torch.float)
b1 = torch.randn(1, requires_grad=True, dtype=torch.float)

model = [Parameter(b0), Parameter(b1)] # Correctly imported Parameter class
```

```

# Set up the optimizer
optimizer = torch.optim.SGD(model, lr=0.1)

# Run optimization loop
for epoch in range(500):
    # Remove the grad computed in the last step
    optimizer.zero_grad()

    # Compute predicted value
    y_predicted = model[0] + model[1] * x_train

    # Calculate the loss function
    # since this is the entire training set, this is effectively using the C
    loss = torch.mean((y_train - y_predicted)**2)

    # Compute gradients
    loss.backward()

    # Update model parameters
    optimizer.step()

    if epoch % 100 == 0:
        print(f"Epoch {epoch}: Loss = {loss.item()}")
        print("Gradient w.r.t b0:", model[0].grad)
        print("Gradient w.r.t b1:", model[1].grad)

# output the final parameters
print("b0:", model[0].item())
print("b1:", model[1].item())

```

```

Epoch 0: Loss = 4.773470401763916
Gradient w.r.t b0: tensor([-2.9494])
Gradient w.r.t b1: tensor([-2.8822])
Epoch 100: Loss = 1.0072836875915527
Gradient w.r.t b0: tensor([-4.8429e-07])
Gradient w.r.t b1: tensor([3.5763e-07])
Epoch 200: Loss = 1.0072836875915527
Gradient w.r.t b0: tensor([-4.8429e-07])
Gradient w.r.t b1: tensor([3.5763e-07])
Epoch 300: Loss = 1.0072836875915527
Gradient w.r.t b0: tensor([-4.8429e-07])
Gradient w.r.t b1: tensor([3.5763e-07])
Epoch 400: Loss = 1.0072836875915527
Gradient w.r.t b0: tensor([-4.8429e-07])
Gradient w.r.t b1: tensor([3.5763e-07])
b0: 1.818988561630249
b1: 1.00831937789917

```